

In [1]:

```
from keras.layers import Input, Conv2D, MaxPooling2D, UpSampling2D, BatchNormalization, Activation
from keras.models import Model
from keras import backend as K

def Encoder():
    input_img = Input(shape=(256, 384, 3)) # adapt this if using `channels_first` image data format
    x = Conv2D(64, (3, 3), padding='same')(input_img)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = MaxPooling2D((2, 2), padding='same')(x)
    x = Conv2D(32, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = MaxPooling2D((2, 2), padding='same')(x)
    x = Conv2D(16, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    encoded = MaxPooling2D((2, 2), padding='same')(x)
    return Model(input_img, encoded)

def Decoder():
    input_img = Input(shape=(32, 48, 16)) # adapt this if using `channels_first` image data format
    x = Conv2D(16, (3, 3), padding='same')(input_img)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(32, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(64, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(3, (3, 3), padding='same')(x)
    x = BatchNormalization()(x)
    decoded = Activation('sigmoid')(x)
    return Model(input_img, decoded)
```

/anaconda3/envs/panicroom/lib/python3.6/site-packages/h5py/__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

from ._conv import register_converters as _register_converters
Using TensorFlow backend.

In [2]:

```
# define input to the model:
x = Input(shape=(256, 384, 3)) # adapt this if using `channels_first` image data format

# make the model:
autoencoder = Model(x, Decoder()(Encoder()(x)))

# compile the model:
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
```

In [3]:

```
Encoder().summary()
```

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	(None, 256, 384, 3)	0
conv2d_8 (Conv2D)	(None, 256, 384, 64)	1792
batch_normalization_8 (Batch Normalization)	(None, 256, 384, 64)	256
activation_8 (Activation)	(None, 256, 384, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 128, 192, 64)	0
conv2d_9 (Conv2D)	(None, 128, 192, 32)	18464
batch_normalization_9 (Batch Normalization)	(None, 128, 192, 32)	128
activation_9 (Activation)	(None, 128, 192, 32)	0
max_pooling2d_5 (MaxPooling2D)	(None, 64, 96, 32)	0
conv2d_10 (Conv2D)	(None, 64, 96, 16)	4624
batch_normalization_10 (Batch Normalization)	(None, 64, 96, 16)	64
activation_10 (Activation)	(None, 64, 96, 16)	0
max_pooling2d_6 (MaxPooling2D)	(None, 32, 48, 16)	0
=====		
Total params: 25,328		
Trainable params: 25,104		
Non-trainable params: 224		

In [4]:

```
Decoder().summary()
```

Layer (type)	Output Shape	Param #
=====		
input_5 (InputLayer)	(None, 32, 48, 16)	0
conv2d_11 (Conv2D)	(None, 32, 48, 16)	2320
batch_normalization_11 (Batch Normalization)	(None, 32, 48, 16)	64
activation_11 (Activation)	(None, 32, 48, 16)	0
up_sampling2d_4 (UpSampling2D)	(None, 64, 96, 16)	0
conv2d_12 (Conv2D)	(None, 64, 96, 32)	4640
batch_normalization_12 (Batch Normalization)	(None, 64, 96, 32)	128
activation_12 (Activation)	(None, 64, 96, 32)	0
up_sampling2d_5 (UpSampling2D)	(None, 128, 192, 32)	0
conv2d_13 (Conv2D)	(None, 128, 192, 64)	18496
batch_normalization_13 (Batch Normalization)	(None, 128, 192, 64)	256
activation_13 (Activation)	(None, 128, 192, 64)	0
up_sampling2d_6 (UpSampling2D)	(None, 256, 384, 64)	0
conv2d_14 (Conv2D)	(None, 256, 384, 3)	1731
batch_normalization_14 (Batch Normalization)	(None, 256, 384, 3)	12
activation_14 (Activation)	(None, 256, 384, 3)	0
=====		
Total params: 27,647		
Trainable params: 27,417		
Non-trainable params: 230		

In [5]:

```
# load all data, normalise, reshape
import os
from PIL import Image
import numpy as np

all_data = []
path = './images-small/'
data_size = 2000
for filename in os.listdir(path)[:data_size]:
    img_path = path + filename
    image = Image.open(img_path)
    all_data.append(np.array(image))
    print('Processed ', len(all_data), ' of ', data_size, ' images.', end='\r'
)

x_train = np.array(all_data[:1500])
x_train = x_train.astype('float32') / 255.
print('x_train shape:', x_train.shape)

x_test = np.array(all_data[1500:])
x_test = x_test.astype('float32') / 255.
print('x_train shape:', x_test.shape)
```

```
x_train shape: (1500, 256, 384, 3)images. of 2000 images.
x_train shape: (500, 256, 384, 3)
```

In [8]:

```
autoencoder.fit(x_train, x_train,
                epochs=25,
                batch_size=64,
                shuffle=True,
                validation_data=(x_test, x_test))
```

Train on 1500 samples, validate on 500 samples

Epoch 1/25

```
1500/1500 [=====] - 1342s 895ms/step - loss: 0.5230 - val_loss: 0.5289
```

Epoch 2/25

```
1500/1500 [=====] - 1296s 864ms/step - loss: 0.5172 - val_loss: 0.5154
```

Epoch 3/25

```
1500/1500 [=====] - 1297s 864ms/step - loss: 0.5128 - val_loss: 0.5138
```

Epoch 4/25

```
1500/1500 [=====] - 1292s 861ms/step - loss: 0.5095 - val_loss: 0.5105
```

Epoch 5/25

```
1500/1500 [=====] - 1305s 870ms/step - loss: 0.5066 - val_loss: 0.5060
```

Epoch 6/25

```
1500/1500 [=====] - 1311s 874ms/step - loss: 0.5038 - val_loss: 0.5106
```

Epoch 7/25

1500/1500 [=====] - 1305s 870ms/step - lo
ss: 0.5014 - val_loss: 0.5139
Epoch 8/25
1500/1500 [=====] - 1294s 863ms/step - lo
ss: 0.4997 - val_loss: 0.5072
Epoch 9/25
1500/1500 [=====] - 1302s 868ms/step - lo
ss: 0.4976 - val_loss: 0.5048
Epoch 10/25
1500/1500 [=====] - 1303s 869ms/step - lo
ss: 0.4964 - val_loss: 0.4963
Epoch 11/25
1500/1500 [=====] - 1303s 869ms/step - lo
ss: 0.4950 - val_loss: 0.5006
Epoch 12/25
1500/1500 [=====] - 1302s 868ms/step - lo
ss: 0.4939 - val_loss: 0.4926
Epoch 13/25
1500/1500 [=====] - 1300s 867ms/step - lo
ss: 0.4925 - val_loss: 0.4970
Epoch 14/25
1500/1500 [=====] - 1300s 867ms/step - lo
ss: 0.4916 - val_loss: 0.4898
Epoch 15/25
1500/1500 [=====] - 1287s 858ms/step - lo
ss: 0.4907 - val_loss: 0.4903
Epoch 16/25
1500/1500 [=====] - 1298s 866ms/step - lo
ss: 0.4899 - val_loss: 0.4929
Epoch 17/25
1500/1500 [=====] - 1310s 873ms/step - lo
ss: 0.4891 - val_loss: 0.4883
Epoch 18/25
1500/1500 [=====] - 1305s 870ms/step - lo
ss: 0.4890 - val_loss: 0.4880
Epoch 19/25
1500/1500 [=====] - 1303s 868ms/step - lo
ss: 0.4880 - val_loss: 0.4872
Epoch 20/25
1500/1500 [=====] - 1298s 865ms/step - lo
ss: 0.4873 - val_loss: 0.4917
Epoch 21/25
1500/1500 [=====] - 1298s 865ms/step - lo
ss: 0.4872 - val_loss: 0.4868
Epoch 22/25
1500/1500 [=====] - 1296s 864ms/step - lo
ss: 0.4865 - val_loss: 0.4876
Epoch 23/25
1500/1500 [=====] - 1292s 861ms/step - lo
ss: 0.4861 - val_loss: 0.4869
Epoch 24/25
1500/1500 [=====] - 1290s 860ms/step - lo
ss: 0.4857 - val_loss: 0.4868
Epoch 25/25
1500/1500 [=====] - 1303s 869ms/step - lo
ss: 0.4855 - val_loss: 0.4856

Out[8]:

<keras.callbacks.History at 0x186cbacdd8>

In [11]:

```
autoencoder.save('./lomo-autoencoder-25epoch-backup')
```

In [12]:

```
# definition to show original image and reconstructed image
def showOrigDec(orig, dec, num=10):
    import matplotlib.pyplot as plt
    %matplotlib inline
    n = num
    plt.figure(figsize=(30, 4))

    for i in range(n):
        # display original
        ax = plt.subplot(2, n, i+1)
        plt.imshow(orig[i].reshape(256, 384, 3))
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)


        # display reconstruction
        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(dec[i].reshape(256, 384, 3))
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
    plt.show()
```

In [15]:

```
encoded_imgs = Encoder().predict(x_train)
decoded_imgs = Decoder().predict(encoded_imgs)
autodecoded_imgs = autoencoder.predict(x_train)
```

In [16]:

```
showOrigDec(x_train, autodecoded_imgs, num=10)
```



In [134]:

```
input_size = np.prod(x_train.shape[1:])
code_size = np.prod(encoded_imgs.shape[1:])
compression_factor = input_size/code_size
print ("Compression factor is ", compression_factor)
```


Compression factor is 12.0

In [30]:

```
import matplotlib.pyplot as plt
%matplotlib inline
plt.imshow(decoded_imgs[1])
```

Out[30]:

<matplotlib.image.AxesImage at 0x186d4afba8>




In [13]:

```
encoded_imgs = Encoder().predict(x_test)
decoded_imgs = Decoder().predict(encoded_imgs)
```

In [31]:

```
showOrigDec(x_test, decoded_imgs, num=10)
```




In [28]:

```
import matplotlib.pyplot as plt
%matplotlib inline
plt.imshow(decoded_imgs[-1])
```

Out[28]:

<matplotlib.image.AxesImage at 0x1870377320>



In []:

```
decoded_imgs[0]
```

In [32]:

```
# load new data for testing the network
import os
from PIL import Image
import numpy as np

new_data = []
path = './images-small/'
for filename in os.listdir(path)[2000:2500]:
    img_path = path + filename
    image = Image.open(img_path)
    new_data.append(np.array(image))
    print('Processed ', len(new_data), ' of ', 500, ' images.', end='\r')

x_val = np.array(new_data[:1500])
x_val = x_val.astype('float32') / 255.
print('x_val shape:', x_val.shape)
```


x_val shape: (500, 256, 384, 3).

In [34]:

```
encoded_imgs = Encoder().predict(x_val)
decoded_imgs = Decoder().predict(encoded_imgs)
autodecoded_imgs = autoencoder.predict(x_val)
```


In [40]:

```
showOrigDec(x_val, autodecoded_imgs, num=10)
```



In [127]:

```
for data in range(0, :  
    img = Image.fromarray((data[0] * 255).astype(np.uint8))  
filename = 1  
img.save('./images-decoded/'+str(filename)+'.jpg')
```

In [107]:

```
scaledup = encoded_imgs*150
```

In [108]:


```
scaledup = Decoder().predict(scaledup[0:1])
```

In [109]:

```
plt.imshow(scaledup[0])
```

Out[109]:

<matplotlib.image.AxesImage at 0x191da30e10>



In [77]:

```
from __future__ import print_function  
from ipywidgets import interact, interactive, fixed, interact_manual  
import ipywidgets as widgets
```

In [85]:

```
def show(scale):  
    scaledup = encoded_imgs*scale  
    scaledup = Decoder().predict(scaledup[0:1])  
    plt.imshow(scaledup[0])  
  
interactive(show, scale=10)
```

